

AVL Trees

Cost of the BST Operations

Our Goal

- Develop a data structure that has **guaranteed** $O(\log n)$ worst-case complexity for **lookup**, **insert** and **find_min**
 - **always!**

	<i>Unsorted array</i>	<i>Array sorted by key</i>	<i>Linked list</i>	<i>Hash Table</i>	
lookup	$O(n)$	$O(\log n)$	$O(n)$	$O(1)$ <small>average and amortized</small>	$O(\log n)$
insert	$O(1)$ <small>amortized</small>	$O(n)$	$O(1)$	$O(1)$ <small>average and amortized</small>	$O(\log n)$
find_min	$O(n)$	$O(1)$	$O(n)$	$O(n)$	$O(\log n)$


BST?

- Do binary search trees achieve this?

Complexity

- Do **lookup**, **insert** and **find_min** have $O(\log n)$ complexity?

- Yes, in this tree



Well, kind of:
we can't talk about
asymptotic complexity
on a single instance

n needs to be
a parameter

- But we are interested in the **worst-case** complexity
- Do **lookup**, **insert** and **find_min** have $O(\log n)$ complexity for *every* BST?


Complexity

- Do **lookup**, **insert** and **find_min** have $O(\log n)$ complexity for *every* BST?

- Consider this sequence of insertions into an initially empty BST

```
insert 10  
insert 20  
insert 30  
insert 40  
insert 50  
insert 60
```

- It produces this tree:



- Then to lookup 70, we have to go through all the nodes
 - This is $O(n)$

This tree has degenerated into a linked list!

- If the insertion sequence is sorted, **lookup** cost $O(n)$

Inserting 70 would also cost $O(n)$

Exercise: find a sequence that yields $O(n)$ cost for **find_min**

Back to Square One


- Develop a data structure that has **guaranteed** $O(\log n)$ worst-case complexity for **lookup**, **insert** and **find_min**
 - **always!**

	<i>Unsorted array</i>	<i>Array sorted by key</i>	<i>Linked list</i>	<i>Hash Table</i>	<i>BST</i>	
lookup	$O(n)$	$O(\log n)$	$O(n)$	$O(1)$ <i>average and amortized</i>	$O(n)$	$O(\log n)$
insert	$O(1)$ <i>amortized</i>	$O(n)$	$O(1)$	$O(1)$ <i>average and amortized</i>	$O(n)$	$O(\log n)$
find_min	$O(n)$	$O(1)$	$O(n)$	$O(n)$	$O(n)$	$O(\log n)$

Something else ...

- BSTs are **not** the data structure we were looking for
 - What else?


Balanced Trees



An Equivalent Tree

- Is there a BST with the same elements that yields $O(\log n)$ cost?


- How about this one?



- It contains the same elements,
- it is sorted,
- but the nodes are arranged differently


Reframing the Problem

- Depending on the tree, BST **lookup** can cost
 - $O(\log n)$ or
 - $O(n)$
- Is there something that remains the same cost-wise?
 - Can we come up with a cost parameter that gives the same complexity in every case?
 - The cost of **lookup** is determined by how far down the tree we need to go
 - if the key is in the tree, the worst case is when it is in a leaf
 - if it is not in the tree, we have to reach a leaf to say so
 - The length of the longest path from the root to a leaf is called the **height** of the tree



Reframing the Problem

- **lookup** for a tree of height h has complexity $O(h)$
 - always!
 - same for **insert** and **find_min**



- But ...
 - h can be in $O(n)$ or in $O(\log n)$
 - where n is the number of nodes in the tree

The Height of a Tree


- The length of the longest path from the root to a leaf
- Let's define it mathematically

$$\left\{ \begin{array}{l} \text{height(\textbf{EMPTY})} = 0 \\ \text{height} \left(\begin{array}{c} \text{ } \\ \swarrow \quad \searrow \\ \triangle_{T_L} \quad \triangle_{T_R} \end{array} \right) = 1 + \max \left(\text{height} \left(\triangle_{T_L} \right), \text{height} \left(\triangle_{T_R} \right) \right) \end{array} \right.$$


This is a recursive definition

Balanced Trees


- A tree is **balanced** if $h \in O(\log n)$
 - where h is its height and n is the number of nodes



Not balanced



Balanced



- On a balanced tree, **lookup**, **insert** and **find_min** cost $O(\log n)$

Self-balancing Trees

New goal:

- make sure that a tree remains balanced as we insert new nodes

... and continues to be a valid BST

● Trees with this property are called **self-balancing**

- There are lots of them

- AVL trees

We will study this one

- Red-black trees

- Splay trees

- B-trees

- ...

Why so many?


- there are many ways to guarantee that the tree remains balanced after each insertion
- some of these tree types have other properties of interest

Self-balancing Trees

- “*the tree stays balanced after each insertion*” is too vague
 - $h \in O(\log n)$ is an asymptotic behavior
 - we can't check it on any given tree
- We want **algorithmically-checkable** constraints that
 1. guarantee that $h \in O(\log n)$
 2. are cheap to maintain
 - at most $O(\log n)$
- We do so by imposing an **additional representation invariants** on trees
 - on top of the ordering invariant
 - this *balance invariant*, when valid, ensures that $h \in O(\log n)$

A Bad Balance Invariant


- Require that
 - (the tree be a BST)
 - all the paths from the root to a leaf have height either h or $h-1$
 - the leaves at height h be on the left-hand side of the tree
- Does it satisfy our requirements?
 1. guarantees that $h \in O(\log n)$ ✓
 - Definitely!
 2. cheap to maintain — at most $O(\log n)$
 - Let's see






A Bad Balance Invariant

- Does it satisfy our requirements?

- 1. guarantees that $h \in O(\log n)$



- Let's insert 5 in this tree



- We changed all the pointers to maintain the balance invariant!

- $O(n)$

- 2. cheap to maintain — at most $O(\log n)$



AVL Trees



Adelson-Velsky

AVL Trees



Landis

The first self-balancing trees (1962)

- **Height invariant**


That's what the balance invariant of AVL trees is called

At every node, the heights of the left and right subtrees differ by at most 1


- An AVL tree satisfies two invariants
 - the ordering invariant
 - the height invariant

The Invariants of AVL Trees

- *The nodes are ordered*
- *At every node, the heights of the left and right subtrees differ by at most 1*
- At any node, there are 3 possibilities




Is this an AVL Tree?



- Is it sorted? ✓
- Do the heights of the two subtrees of every node differ by at most 1? ✓

YES


Is this an AVL Tree?



- Is it sorted? ✓
- Do the heights of the two subtrees of every node differ by at most 1? ✓

YES

Is this an AVL Tree?




- Is it sorted? ✓
- Do the heights of the two subtrees of every node differ by at most 1? ✗
 - It doesn't hold at node 15

- We say there is a **violation** at node 15

NO

Is this an AVL Tree?




- Is it sorted?
- Do the heights of the two subtrees of every node differ by at most 1?



YES

Is this an AVL Tree?




- Is it sorted?
- Do the heights of the two subtrees of every node differ by at most 1?
 - There is a **violation** at node 15
 - and another **violation** at node 10



NO

Is this an AVL Tree?



- Is it sorted?
- Do the heights of the two subtrees of every node differ by at most 1?



The height invariant does **not** imply that the length of every path from the root to a leaf differ by at most 1

YES

Rotations

Insertion Strategy

1. Insert the new node as in a BST

- this preserves the ordering invariant
- but it **may break the height invariant**

2. Fix any height invariant violation

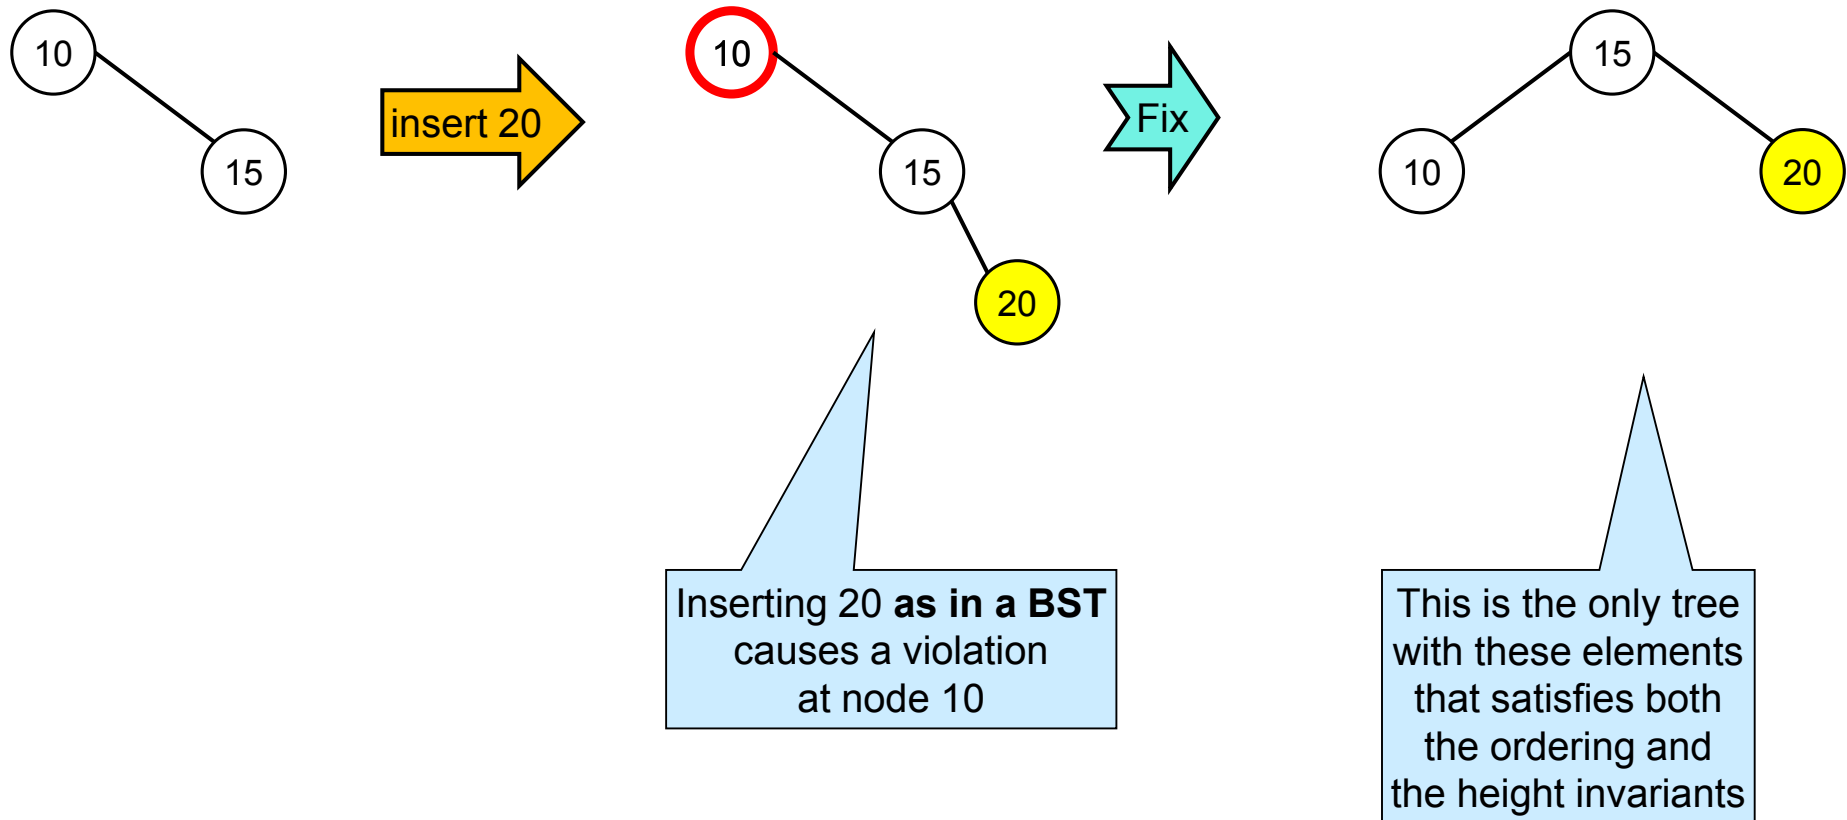
- fix the **lowest** violation
 - this will take care of all other violations

We will see why later




● This is a common approach

- of two invariants, preserve one and temporarily break the other
- then, patch the broken invariant
 - cheaply

Example 1



Example 2



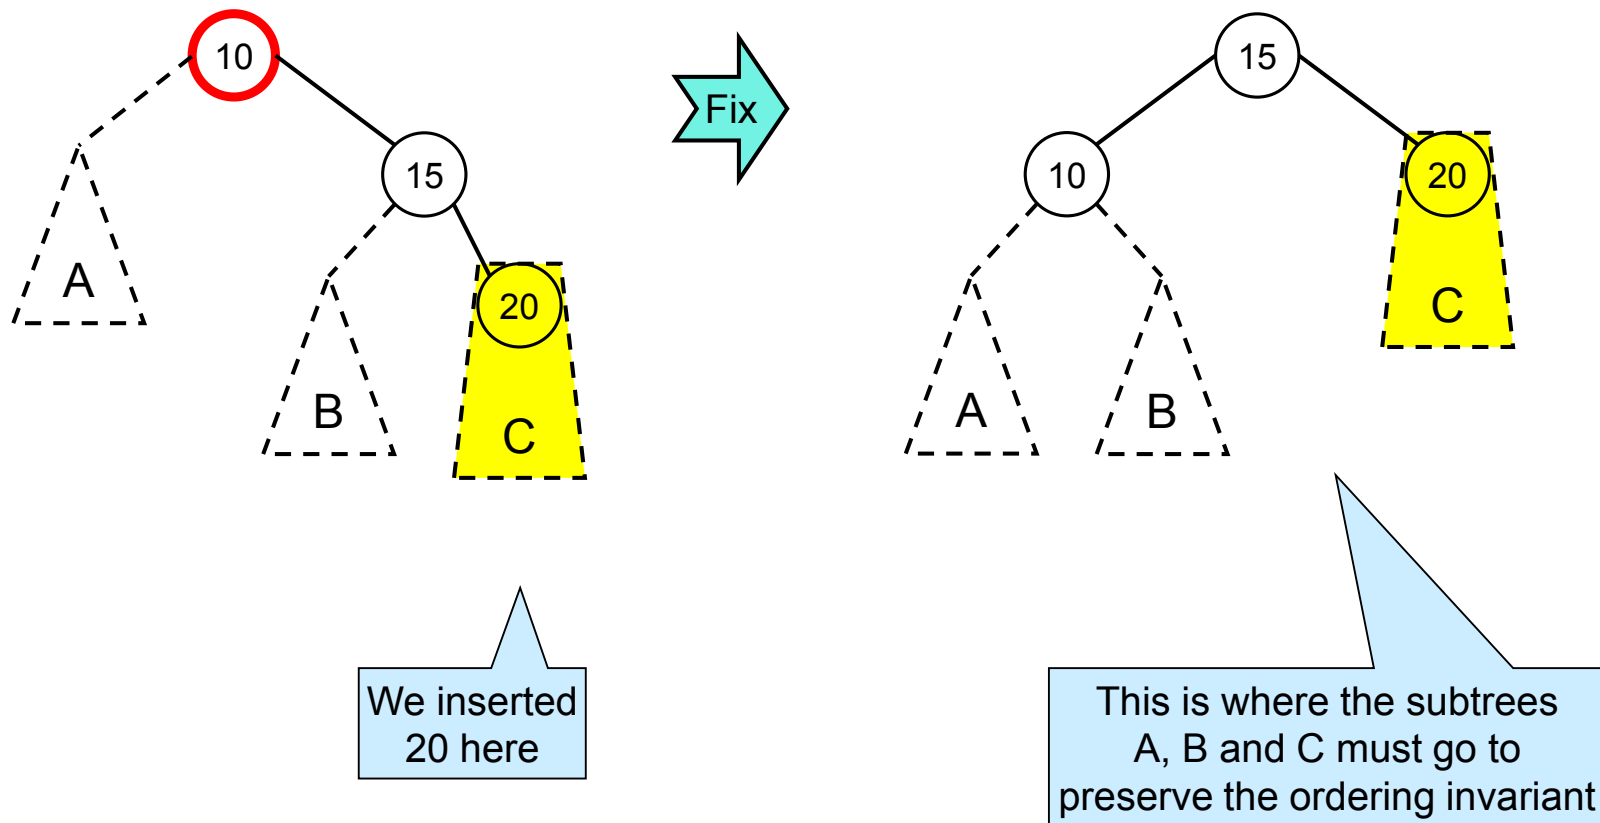
?

Inserting 25 as in a BST causes a violation at node 10

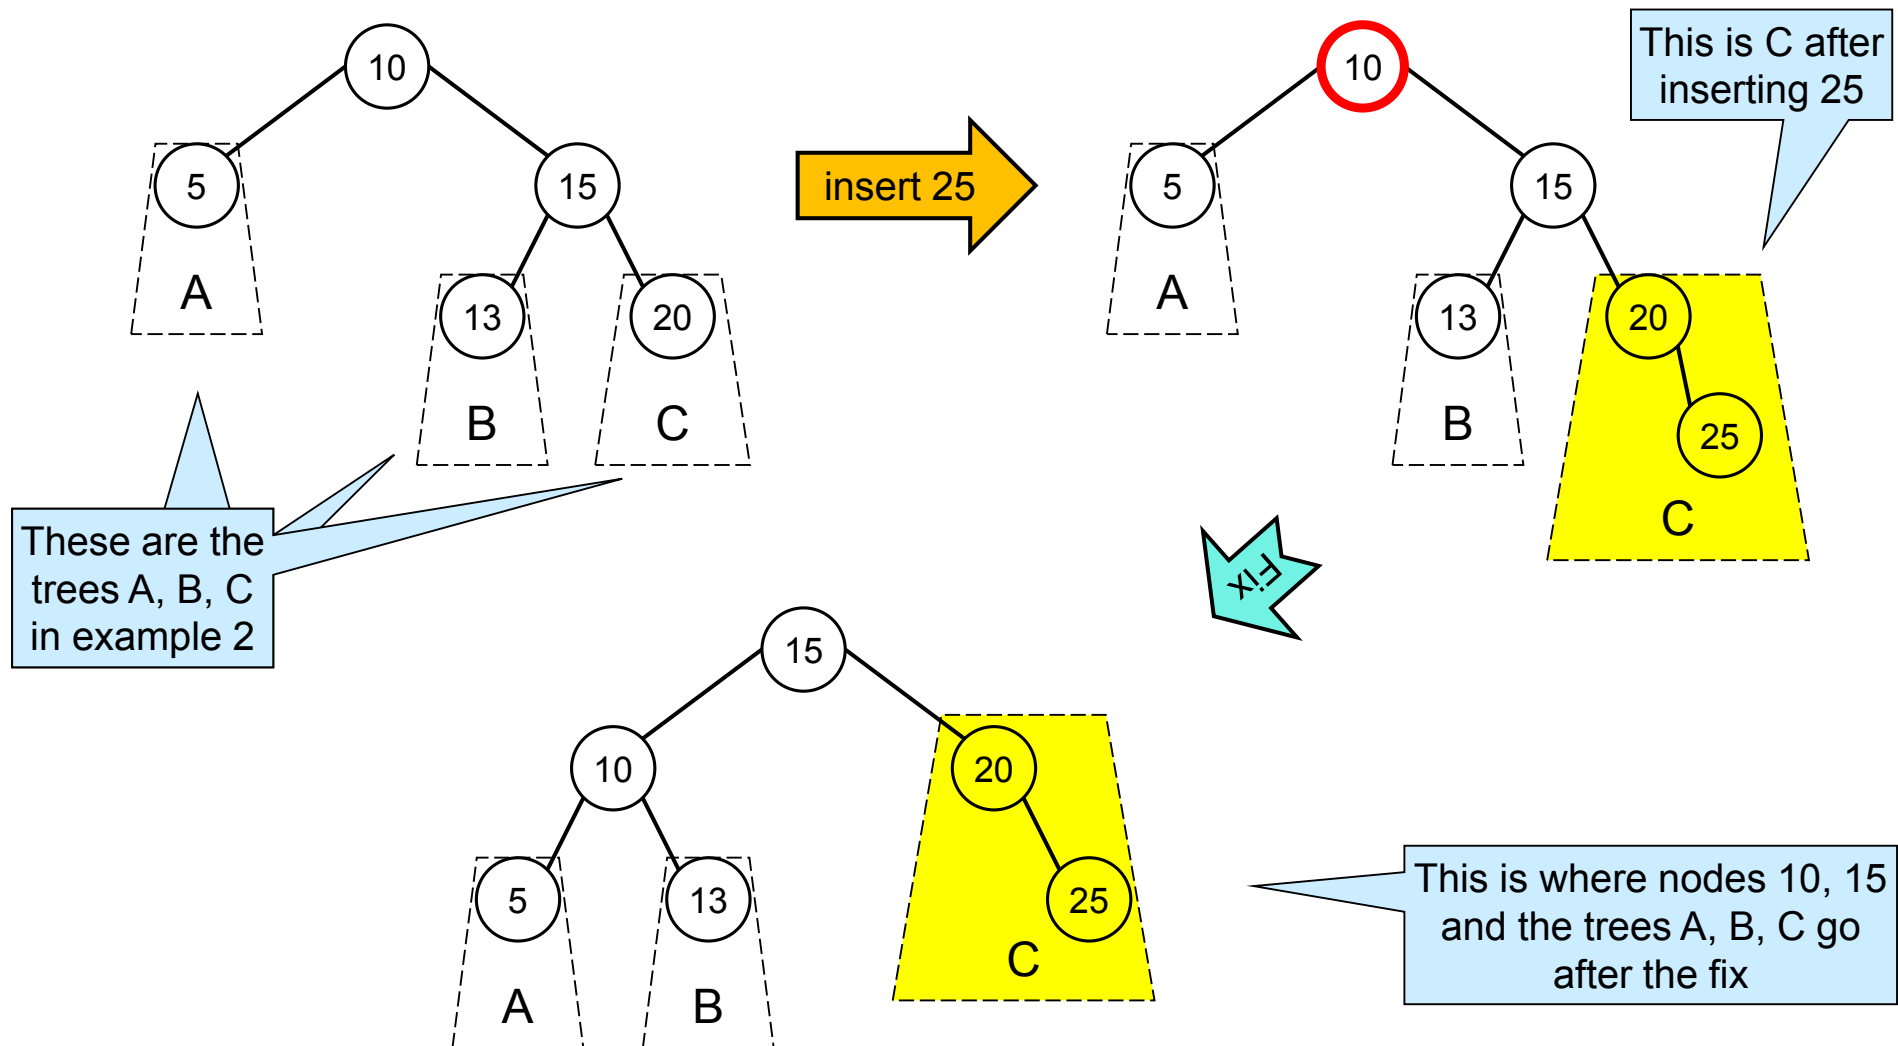
There are a lot of AVL trees with these elements:
which one to pick?

Example 1 Revisited

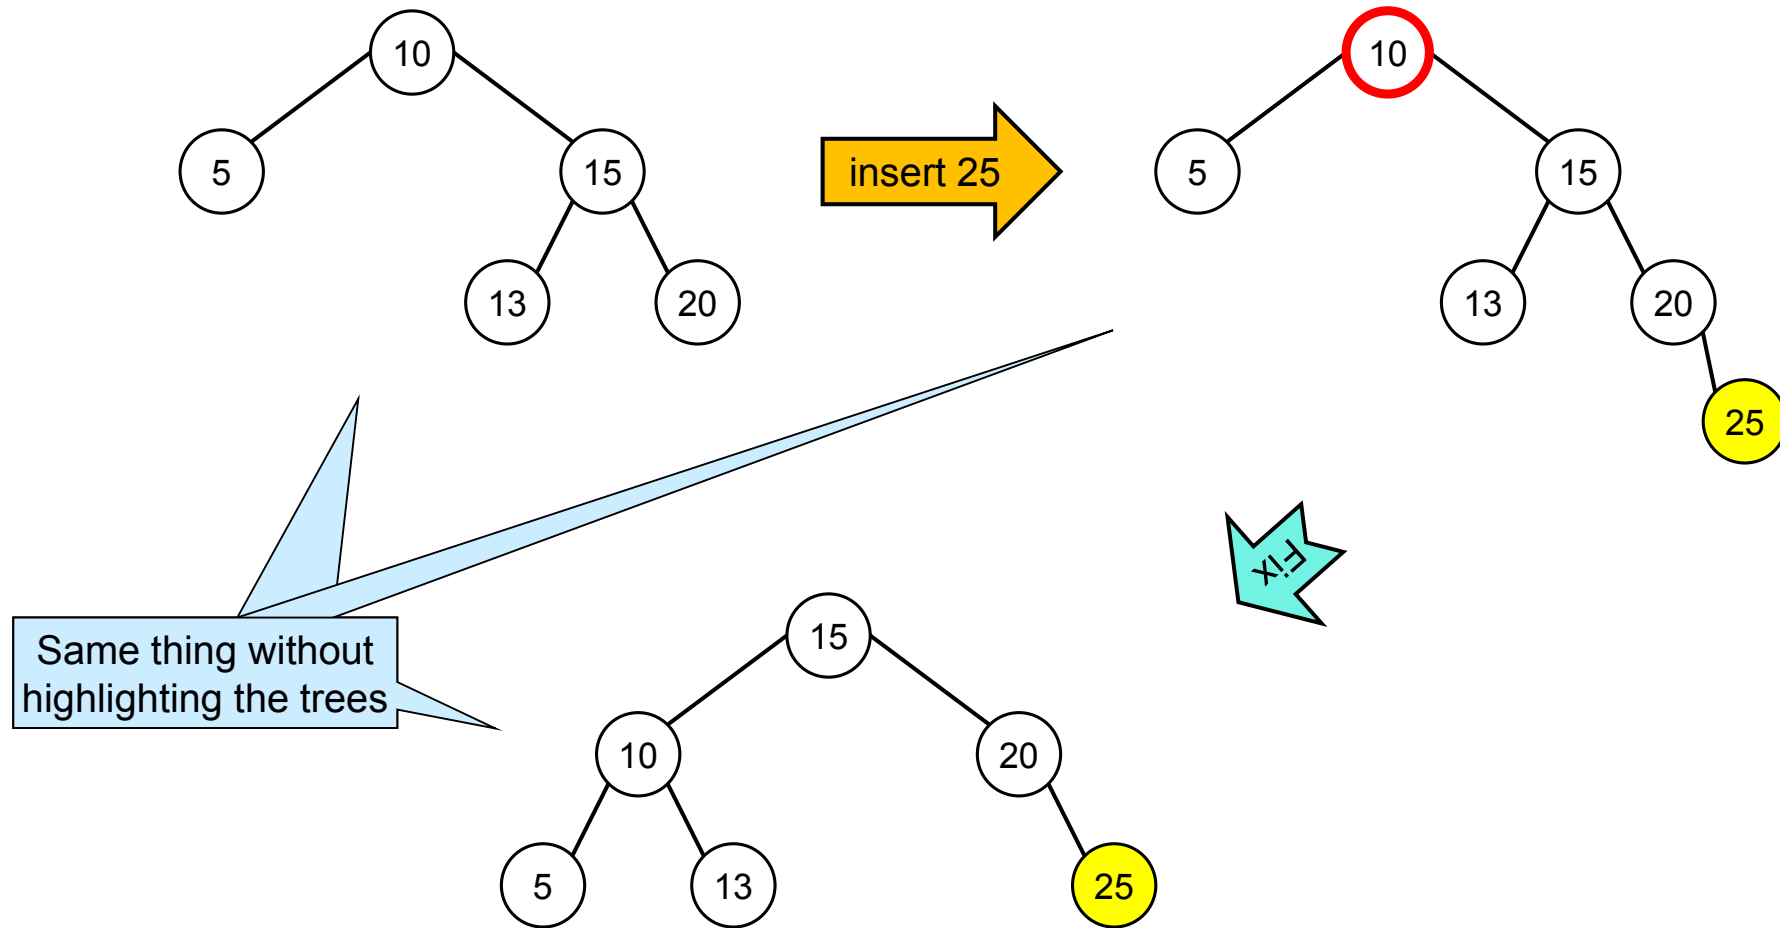
- If this example was part of a bigger tree, what would it look like?



Example 2




Example 2



Left Rotation


- This transformation is called a **left rotation**



- Note that it maintains the ordering invariant
- We do a left rotation when C has become too tall after an insertion

Right Rotation


- The symmetric situation is called a **right rotation**



- It too maintains the ordering invariant
- We do a right rotation when A has become too tall after an insertion

Single Rotations Summary

- Right and left rotations are **single rotations**

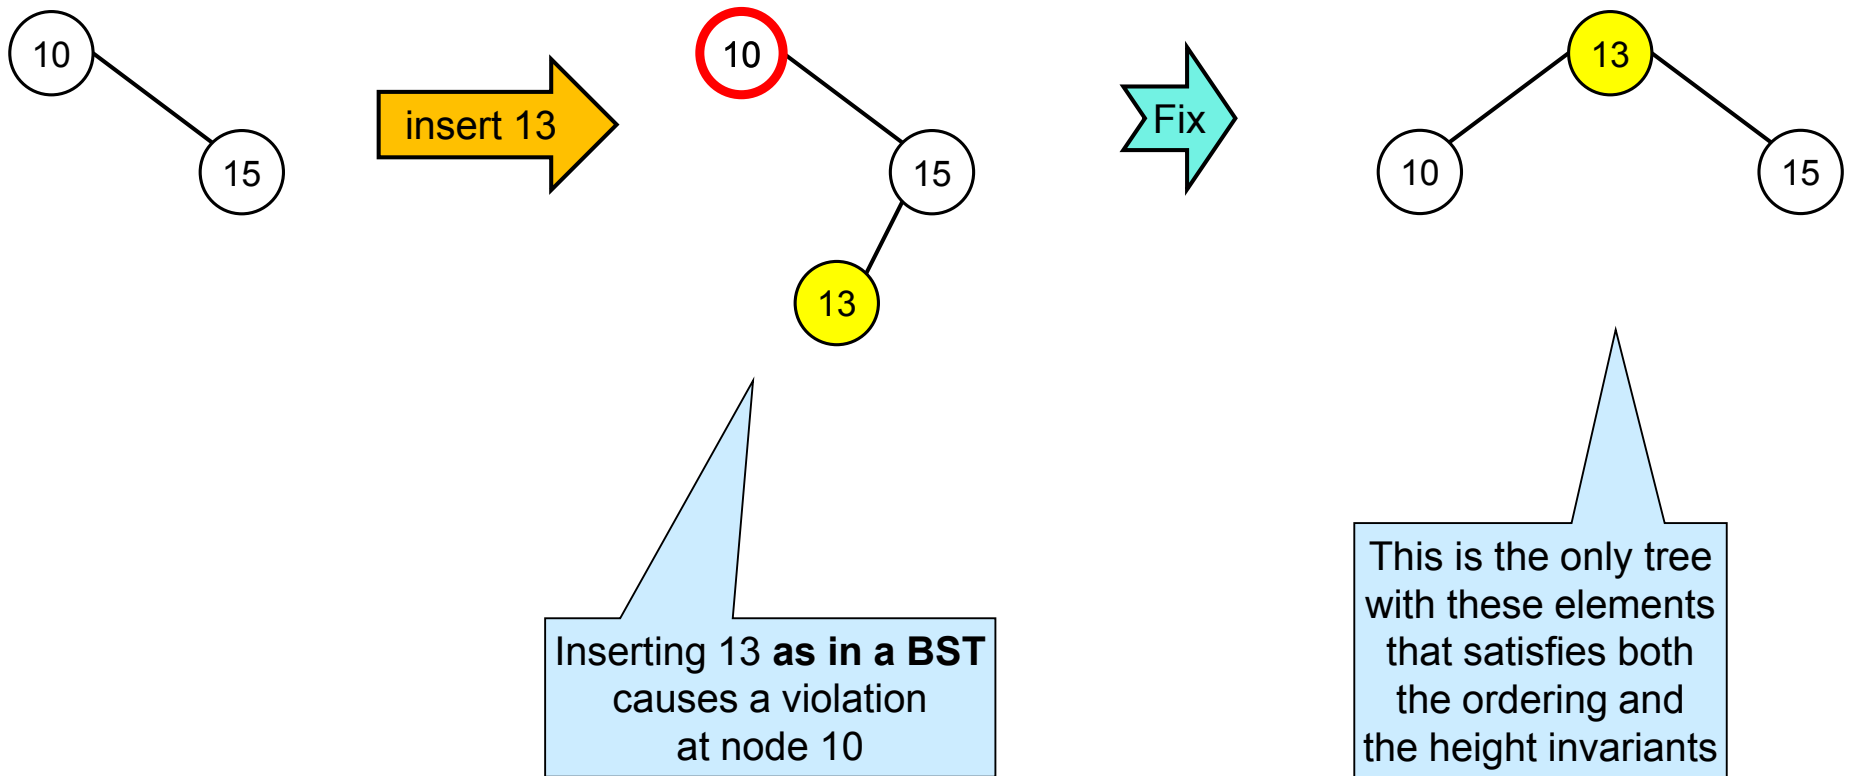


- They maintain the ordering invariant
- We do one of them when
 - the **lowest violation** is at the root
 - one of the **outer subtrees** has become too tall

That's either y or x


That's either A or C respectively

Example 3



Double Rotations

- We can generalize this example to the case where the nodes have subtrees




- This is called a **double rotation**
 - specifically a right-left double rotation

This is where the subtrees A, B, C and D must go to preserve the ordering invariant

Right-left Double Rotation

- Here's the general pattern




- We do this double rotation when the subtree rooted at y has become too tall after an insertion

The ordering invariant is maintained

Left-right Double Rotation


- The symmetric transformation is a left-right double rotation



- We do this double rotation when the subtree rooted at y has become too tall after an insertion

The ordering invariant is maintained

Double Rotations Summary

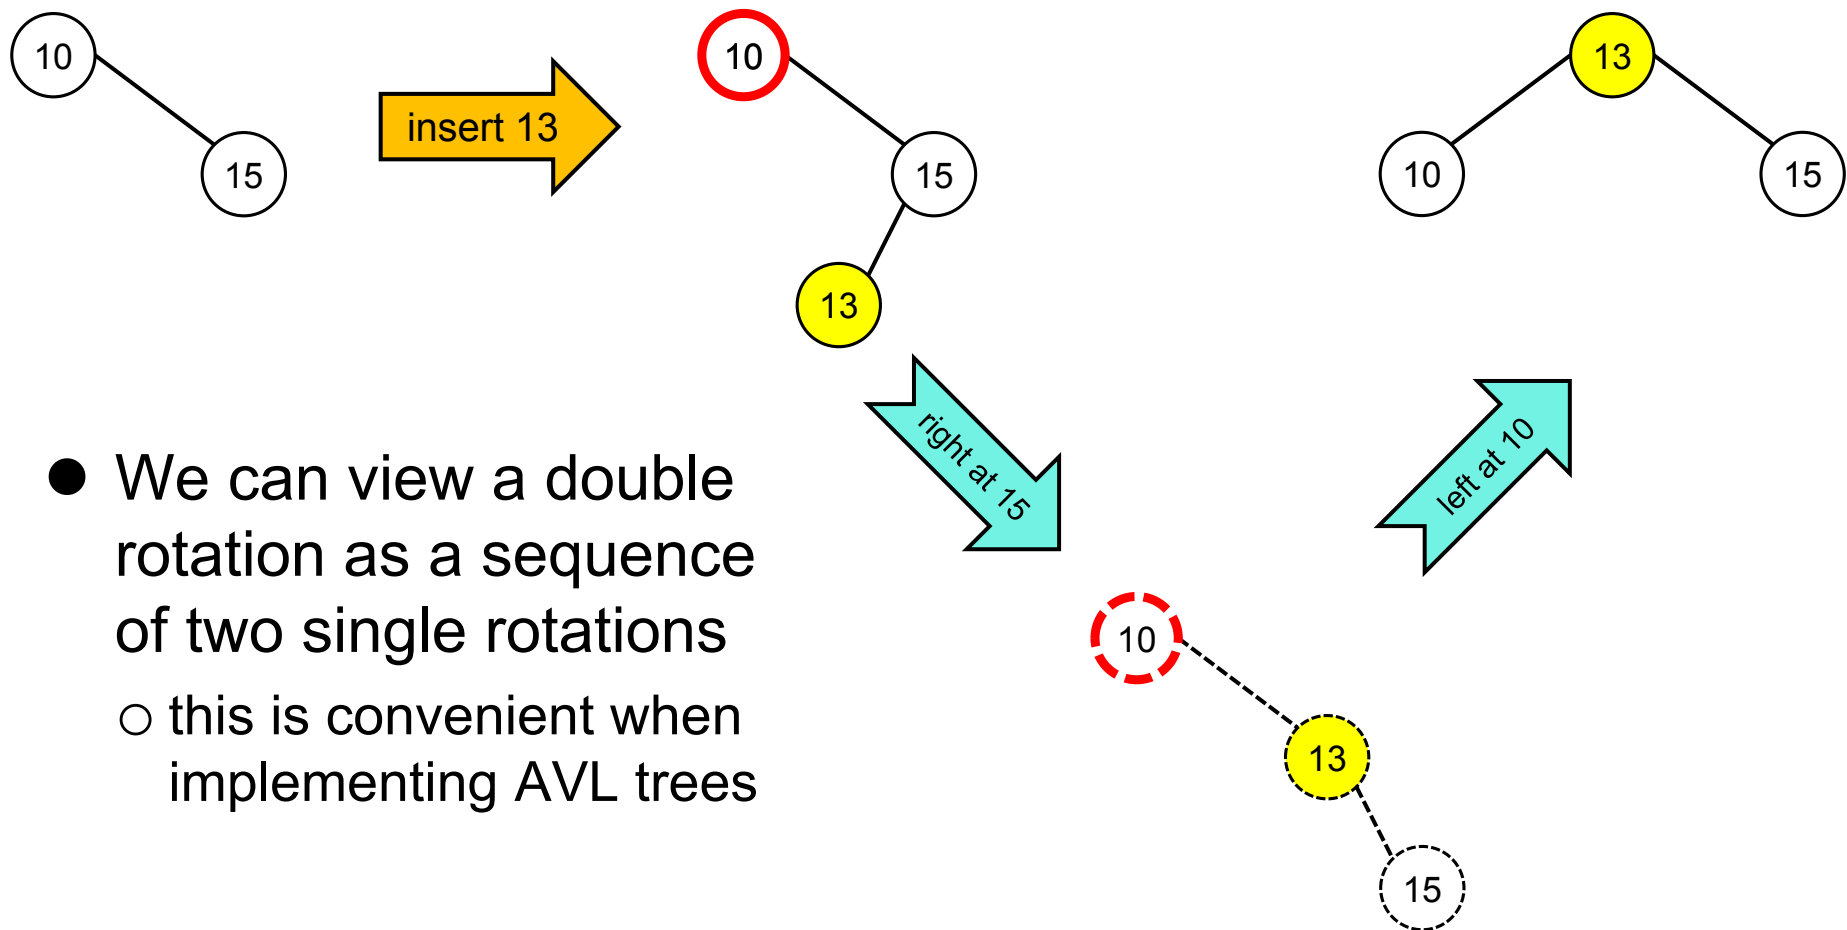


- Double rotations maintain the ordering invariant
- We do one of them when
 - the **lowest violation** is at the root
 - one of the **inner subtrees** has become too tall

That's either z or x



That's the subtree rooted at y

Why is it Called a *Double* Rotation?




AVL Rotation When-to


If the insertion
that caused the lowest
violation **x** happened ...




**right
single
rotation
at x**



**left/right
double
rotation
at x**



**right/left
double
rotation
at x**



**left
single
rotation
at x**

... then do a ...

Self-balancing Requirements

- Does the height constraint satisfy our requirements?

1. It guarantees that $h \in O(\log n)$



Left as exercise

2. It is cheap to maintain — at most $O(\log n)$

- each type of rotation costs $O(1)$

- at most one rotation is needed for each insertion

We will see why next

So, maintaining the height invariant costs $O(1)$



Height Analysis

Insertion into an AVL Tree

- Assume we are inserting a node into an AVL tree of height h

One of two things can happen:


1. This causes a height **violation**

- we fix it with a rotation
 - the resulting tree is a valid AVL tree
- the fixed tree still has height h
 - the tree does not grow

Let's
see
why


2. This does not cause a violation

- the resulting tree has height h or $h+1$
 - the tree may grow only when there is no violation




Fixing the Lowest Violation

- Assume an insertion causes a **violation**
 - possibly more than one




- We will focus on the subtree under the **lowest violation**
 - We will find that fixing it yields a subtree with the **same height h** as the original subtree
 - This necessarily resolves all violations above it
 - because the height of this subtree has not changed
 - if it satisfied the height invariant for the nodes above it before, it still satisfies it after




Fixing the lowest violation fixes the whole tree

The Lowest Violation




- Let's expand the tree
 - T cannot be empty — No violation possible
 - the new node can have been inserted in its left or right subtree
- Let's consider insertion in T_R — Insertion in T_L is symmetric



- To have a violation
 - T_R must be taller than T_L
 - $h-1$ vs. $h-2$
 - T_R must have grown after the insertion
 - from $h-1$ to h

The right subtree has become **too tall**

The Lowest Violation






- Let's expand the right subtree

- T_R cannot be empty


No violation possible

- the new node can have been inserted in its left or right subtree




- *Let's examine each case in turn*

Insertion in the Outer Subtree



- How tall are T_i and T_o ?



- $h_o = h-2$

- T_o needs to be as tall as possible to causes the violation


- $h_i = h_o = h-2$

- h_i may be either $h-2$ or $h-3$





- but if h_i were $h-3$, the lowest violation would be here

T_i and T_o have the same height

Insertion in the Outer Subtree




- T_i and T_o have height $h-2$




- This is the situation where we do a **single left rotation**

○ *Is this an AVL tree?*

Insertion in the Outer Subtree



● Is this an AVL tree?



○ BST insertion and the rotations maintains the **ordering invariant**

○ T_L , T_i and T'_o are AVL trees

➤ because x was the lowest violation

○ T_L-x-T_i is an AVL tree of height $h-1$

➤ because both T_L and T_i have height $h-2$


○ $(T_L-x-T_i)-y-T'_o$ is an AVL tree of height h

➤ because T'_o also has height $h-1$


The **height invariant**
is restored



Insertion in the Inner Subtree




- How tall are T_i and T_o ?




- $h_i = h-2$
 - T_i needs to be as tall as possible to causes the violation
- $h_o = h_i = h-2$
 - h_o may be either $h-2$ or $h-3$
 - but if h_o were $h-3$, the lowest violation would be **here**

T_i and T_o have the same height

Insertion in the Inner Subtree



- T_i and T_o have height $h-2$



- T'_i contains at least the inserted node

➤ let's expand it

- T_1 and T_2 have height $h-2$ or $h-3$


➤ one of them has height $h-2$


- the inserted node could be


➤ the root – if T_1 and T_2 are empty


➤ in T_1

➤ in T_2






- 




53

Insertion in the Inner Subtree




- Is this an AVL tree?




- BST insertion and the rotations maintains the **ordering invariant**

Insertion in the Inner Subtree



● Is this an AVL tree?



- T_L , T_1 , T_2 and T_0 are AVL trees
 - because x was the lowest violation
- T_L-x-T_1 is an AVL tree of height $h-1$
 - because T_L has height $h-2$ and
 - T_1 has height either $h-2$ or $h-3$
- T_2-z-T_0 is an AVL tree of height $h-1$
 - because T_2 has height either $h-2$ or $h-3$
 - T_0 has height $h-2$ and
- $(T_L-x-T_1)-y-(T_2-z-T_0)$ is an AVL tree of height h

The height invariant is restored



Summary

- When inserting into an AVL tree of height h
 - If there is no violation, the tree height remains h or grows to $h+1$
 - If there is a violation, the tree height remains h
- To fix a violation
 - perform a rotation on the **lowest violation**
 - a **single rotation** if the node was inserted in its **outer subtree**
 - a **double rotation** if the node was inserted in its **inner subtree**
- **One** rotation fixes the whole tree
 - The resulting tree is again an AVL tree
 - **lookup**, **insert** and **find_min** cost $O(\log n)$ in it
 - where n is the number of nodes