

Searching Sorted Data

Searching for a Number

- Consider the following sorted array

0	1	2	3	4	5	6	7	8	9
-2	0	4	7	12	19	22	42	65	

- When searching for a number x using binary search, we **always** start by looking at the midpoint, index 4

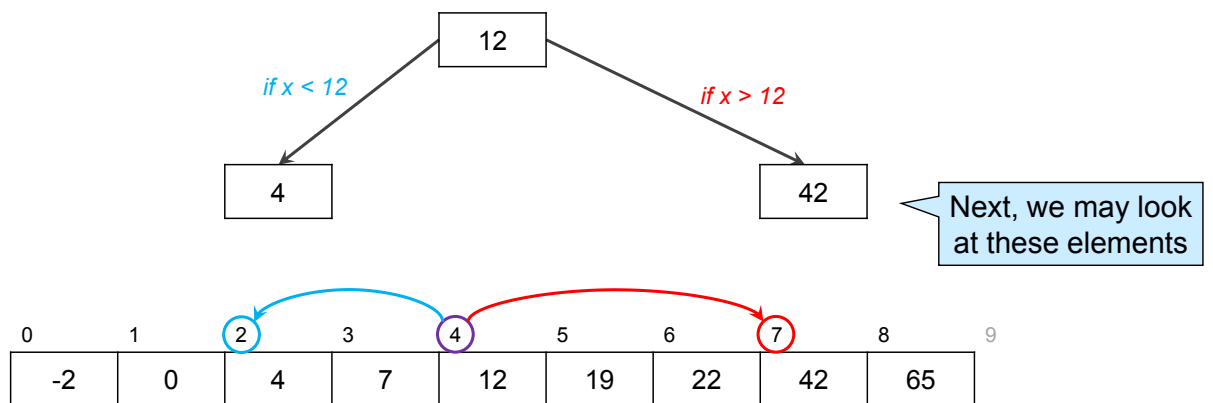
12

We always look at this element

- Then, 3 things can happen
 - $x = 12$ (and we are done)
 - $x < 12$
 - $x > 12$


Searching for a Number

- If $x < 12$, the next index we look at is **necessarily** 2
- If $x > 12$, the next index we look at is **necessarily** 7



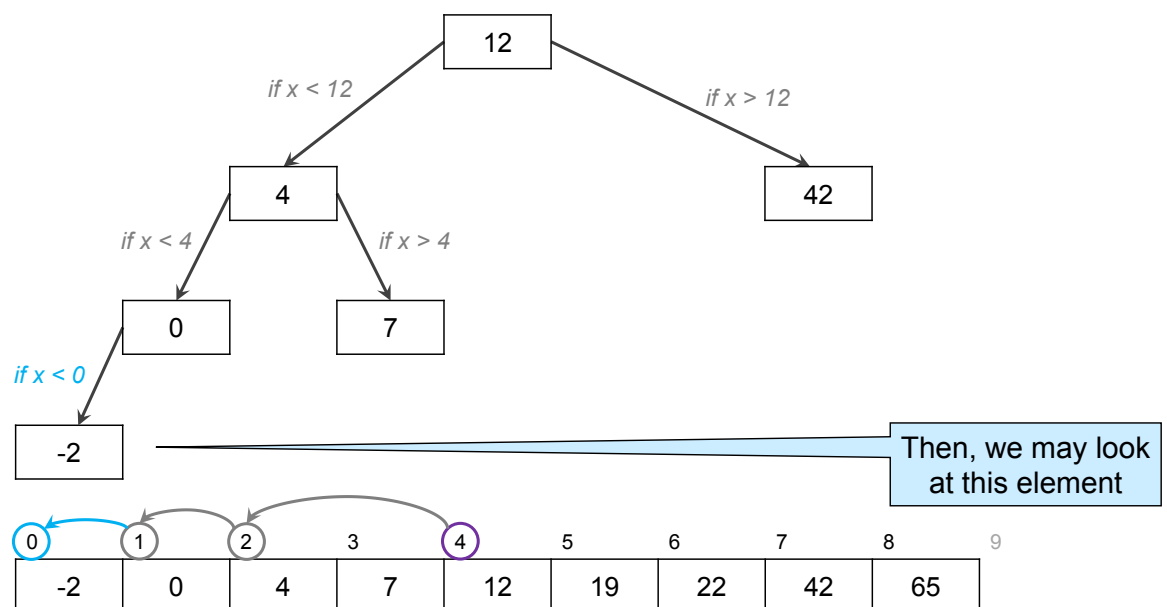
Searching for a Number

- Assume $x < 12$, so we look at 4
 - if $x = 4$, we are done
 - if $x < 4$, we **necessarily** look at 0
 - if $x > 4$, we **necessarily** look at 7



Searching for a Number


- Assume $x < 4$, so we look at 0
 - if $x = 0$, we are done
 - if $x < 0$, we **necessarily** look at -2



Searching for a Number

- We can map out all possible sequences of elements binary search may examine, for any x

This is called a **decision tree**: at every step, it tells us how to decide what to do next

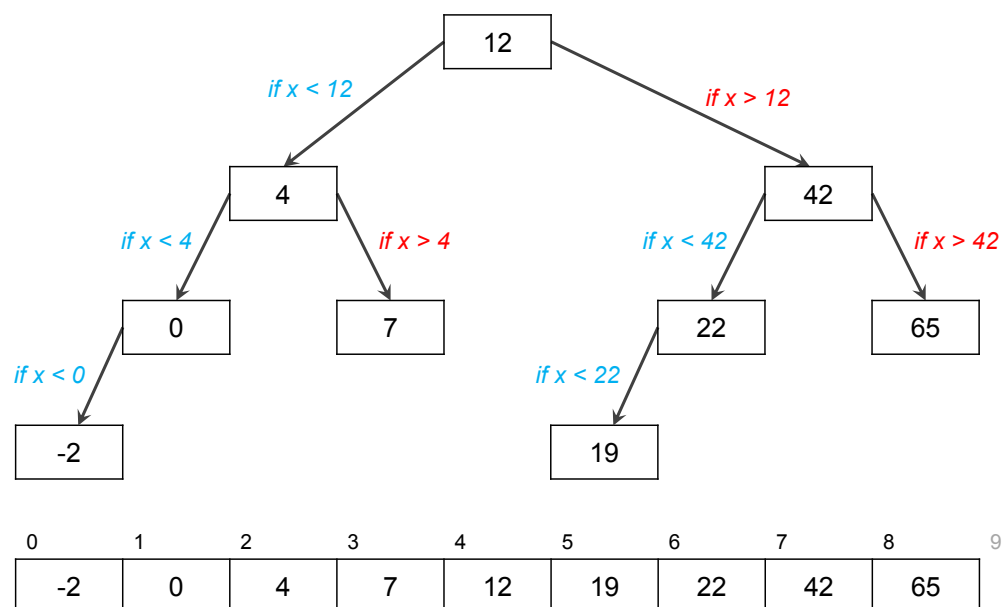


We are essentially hoisting the array by its midpoint, its two sides by their midpoint, etc

0	1	2	3	4	5	6	7	8	9
-2	0	4	7	12	19	22	42	65	

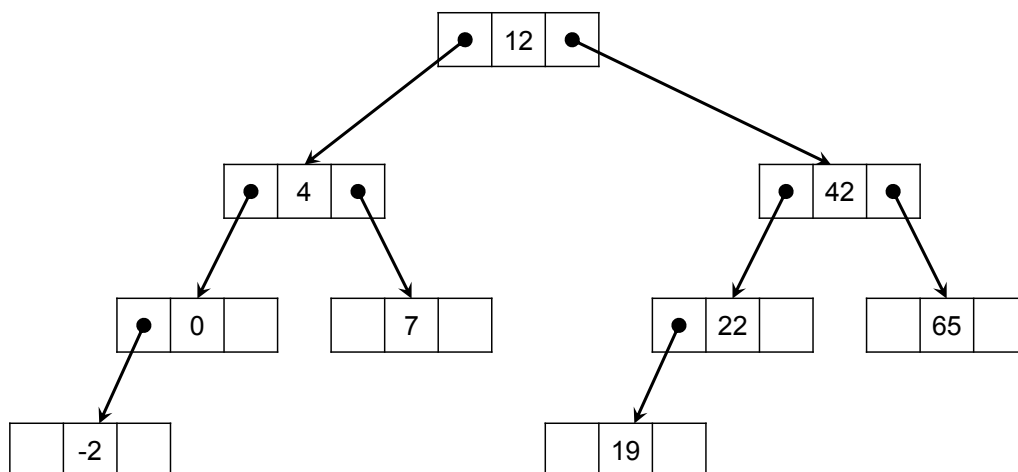
Searching for a Number

- An array provides direct access to all elements
 - This is overkill for binary search
 - At any point, it needs direct access to **at most two** elements



Searching for a Number

- We can achieve the same access pattern by pairing up each element with **two pointers**
 - one to each of the two elements that may be examined next



- We are losing direct access to arbitrary elements,
 - but it retains access to the elements that matter to binary search

Arrays gave us more power than needed

Towards an Implementation

- We can capture this idea with this **type declaration**:

```
typedef struct tree_node tree;  
struct tree_node {  
    tree* left;  
    int data;  
    tree* right;  
};
```

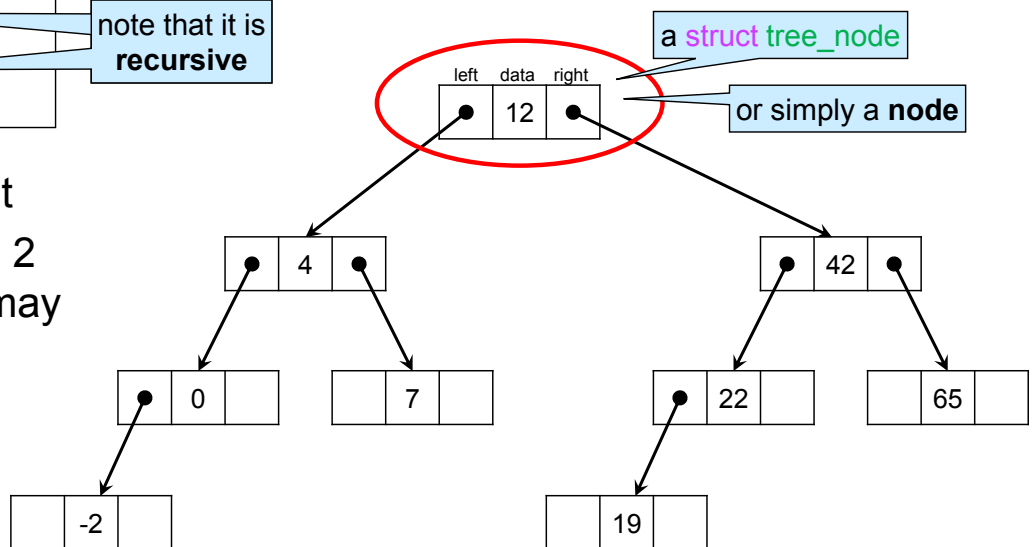
note that it is
recursive

a **struct tree_node**

or simply a **node**

- a data element
- pointers to the 2 elements we may look at next

- This is called a **node**



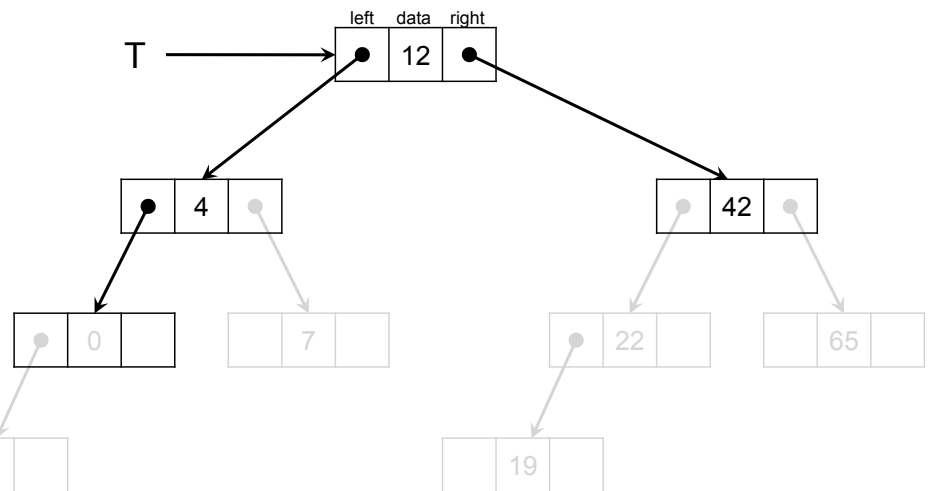
- This arrangement of data in memory is called a **tree**

Constructing this Tree

```
typedef struct tree_node tree;
struct tree_node {
    tree* left;
    int data;
    tree* right;
};
```

- Let's build the first few nodes of this example

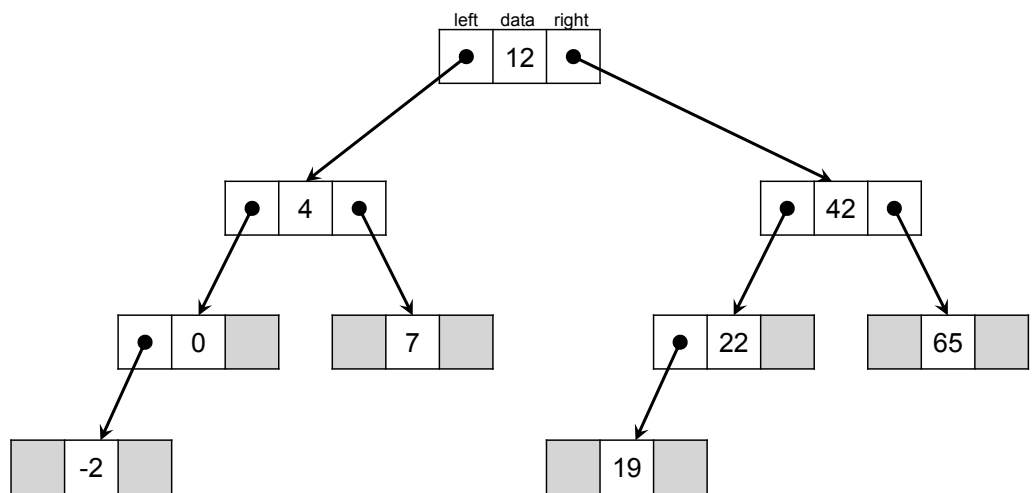
```
tree* T = alloc(tree);
T->data = 12;
T->left = alloc(tree);
T->left->data = 4;
T->right = alloc(tree);
T->right->data = 42;
T->left->left = alloc(tree);
...
```



The End of the Line

```
typedef struct tree_node tree;
struct tree_node {
    tree* left;
    int data;
    tree* right;
};
```

- What should the blank left/right fields point to?



- NULL



➤ each sequence of left/right pointers works like a NULL-terminated list

- a dummy node



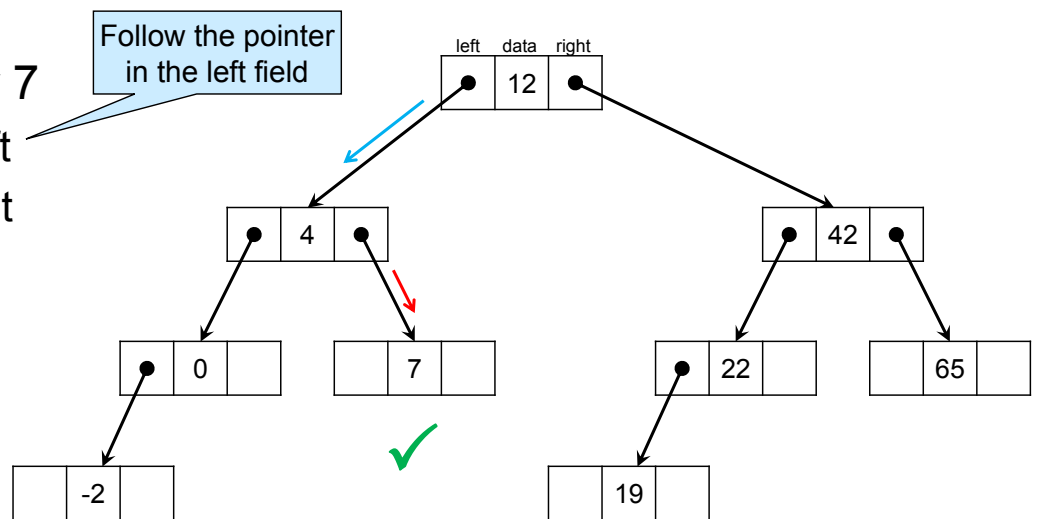
➤ not very useful here

We used dummy nodes to get direct access to the end of a list

Searching

- Searching for 7

- $7 < 12$: go left
- $7 > 4$: go right
- $7 = 7$: **found**




- We are doing the **same steps** as binary search
- Starting from an n-element array, the cost is $O(\log n)$

If the tree is obtained as in this example

Searching

- Searching for 5

- $5 < 12$: go left
- $5 > 4$: go right
- $5 < 7$: go left
 - nowhere to go
- **not there**



- We are doing the **same steps** as binary search
- Starting from an n -element array, the cost is $O(\log n)$

If the tree is obtained
as in this example

Recall our Goal

- Develop a data structure that has **guaranteed** $O(\log n)$ worst-case complexity for **lookup**, **insert** and **find_min**
 - always!

- **lookup** has cost $O(\log n)$

in *this* setup

	<i>Target data structure</i>
lookup	$O(\log n)$
insert	$O(\log n)$
find_min	$O(\log n)$




- What about **insert** and **find_min**?

Insertion

- Inserting 5

- $5 < 12$: go left
- $5 > 4$: go right
- $5 < 7$: go left
- put it there



- We are doing the **same steps** we would do to search for it, and then put it where it should have been
 - so that we find it when searching for it next time
- For an n-element array, this costs $O(\log n)$

We couldn't get this with sorted arrays

If the tree is obtained as in this example

Finding the Smallest Key

- Keep going left


- left from 12 to 4

- left from 4 to 0

- left from 0 to -2

- nothing to its left

- **the smallest key is -2**



- Starting from an n-element array, we can go left at most $O(\log n)$ times

- The cost is $O(\log n)$

If the tree is obtained as in this example

Recall our Goal

- Develop a data structure that has **guaranteed** $O(\log n)$ worst-case complexity for **lookup**, **insert** and **find_min**
 - always!

- **lookup**, **insert** and **find_min** all have cost $O(\log n)$

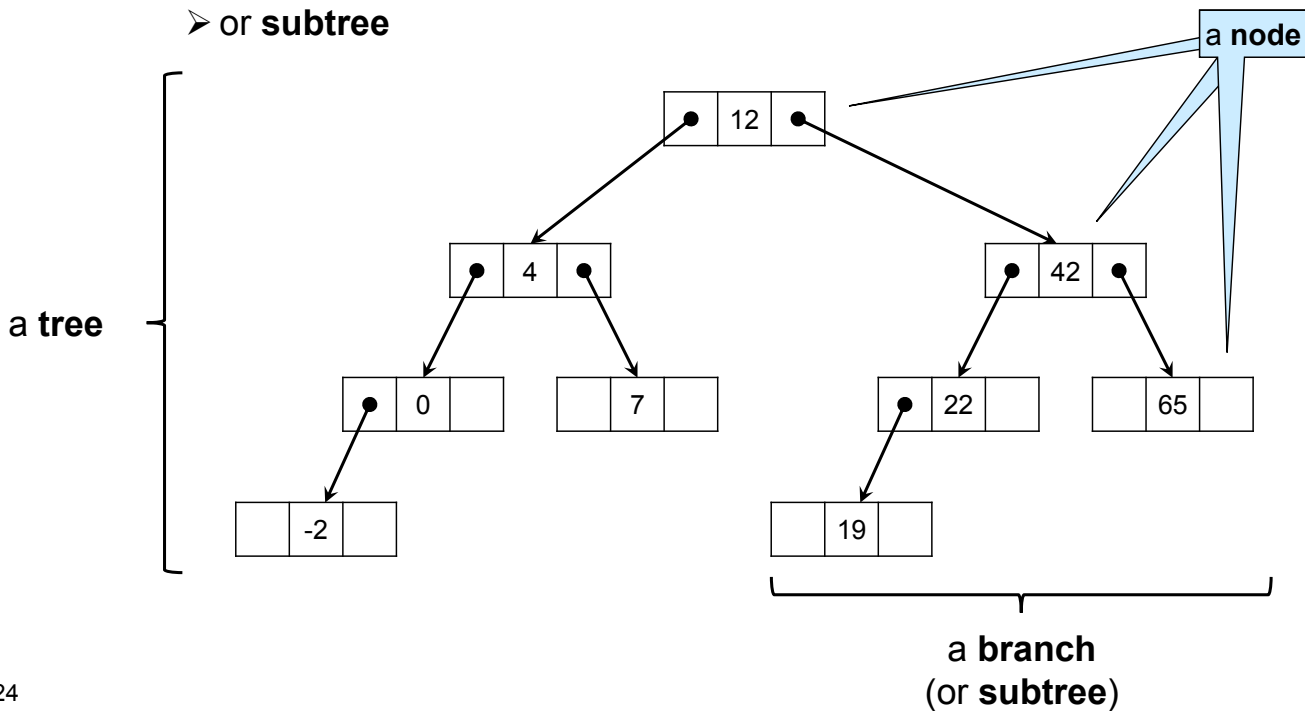
in *this* setup

	<i>Target data structure</i>	
lookup	$O(\log n)$	✓
insert	$O(\log n)$	✓
find_min	$O(\log n)$	✓

Trees

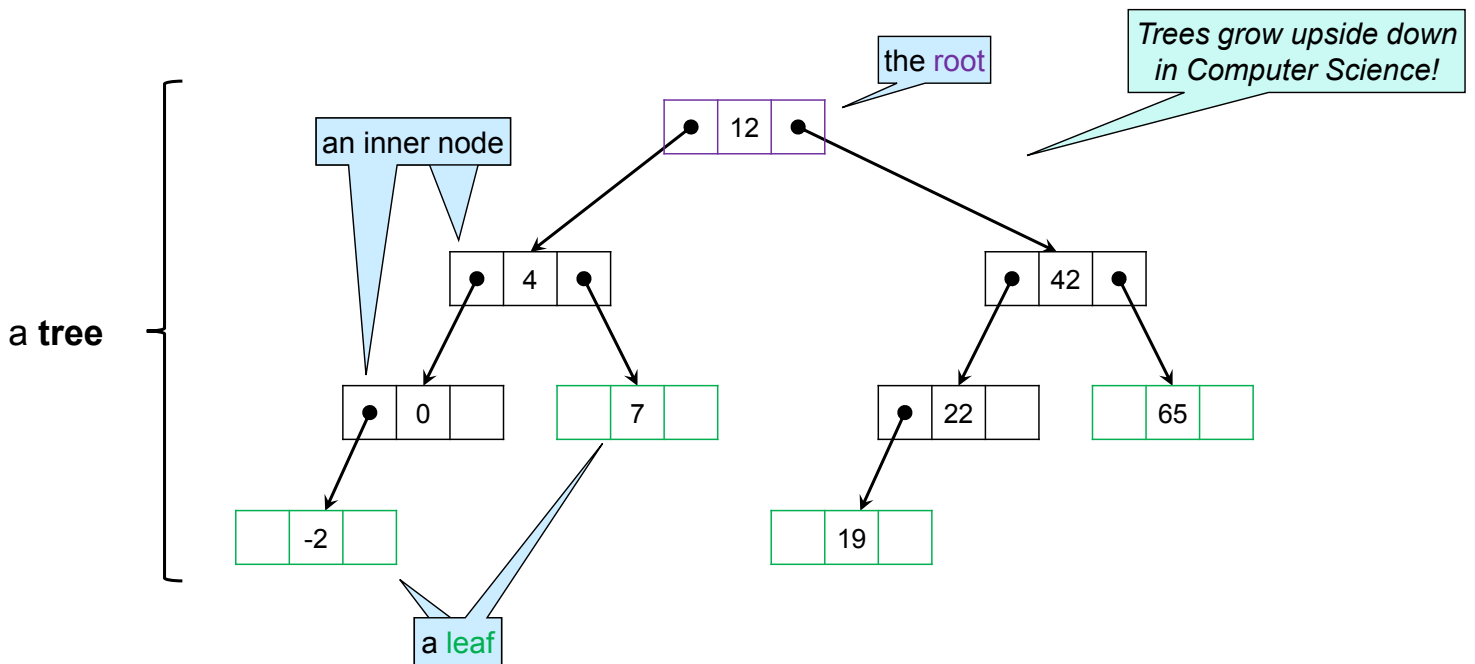
Terminology

- This arrangement of data is called a (binary) **tree**
 - each item in it is called a **node**
 - the part of a tree hanging from a node is called a **branch**
 - or **subtree**



Terminology

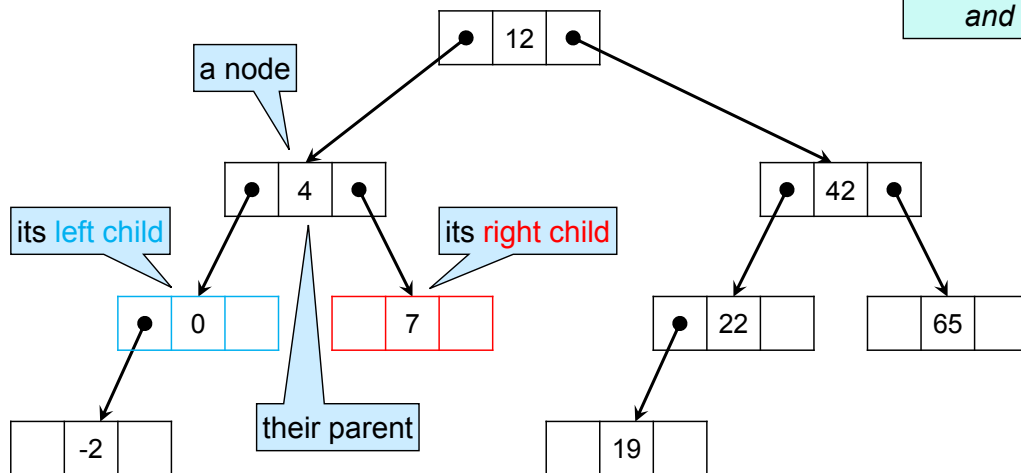
- The node at the top is called the **root** of the tree
 - the nodes at the bottom are the **leaves** of the tree
 - the other nodes are called **inner nodes**



Terminology


- Given any node
 - the node to its left is its **left child**
 - the node to its right is its **right child**
 - the node above it is its **parent**

.. and Computer Science mixes botanical trees and family trees!



Concrete Tree Diagrams


- When drawing trees, we generally omit the details of the memory diagrams
 - draw just the data in a node
 - not the pointer fields
 - and the connection to its children
 - we always draw the root at the top



Pictorial Abstraction

- We will often reason about trees that are arbitrary
 - their actual content is unimportant, so we abstract it away


- We draw a generic tree as a triangle




- We represent the empty tree by simply writing "Empty"

Empty

What do Trees Look Like?




- Abstract trees come in many shapes



- When working with trees, we need to account for all these possibilities
 - we will forget some
- *Is there a simpler description?*

What Trees Look Like


- A tree can be



- either empty





EMPTY

- or a root with a tree on its left and a tree on its right



- **Every tree** reduces to these two cases

EMPTY



the empty tree

a tree with just the root


a root and two children

a tree with just the left subtree

a tree with just the right subtree

What Trees Look Like


- A tree can be



- either empty

- or a root with a tree on its left and a tree on its right

EMPTY



- We only need to consider these two cases when

- writing code about trees
 - reasoning about trees

A Minimal Tree Invariant

- We only need to consider these two cases when writing code about trees

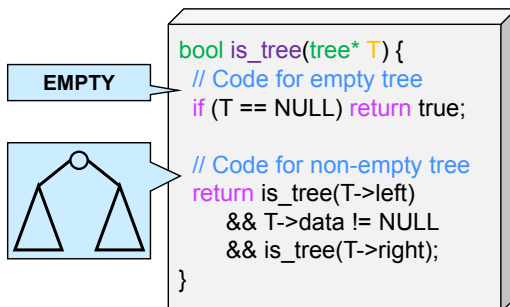
- Let's apply this to write a basic invariant about trees of *entries*

- Just check that the data field is never NULL

```
typedef struct tree_node tree;  
struct tree_node {  
    tree* left;  
    entry data; // != NULL  
    tree* right;  
};
```

Recall we are using trees to implement dictionaries:

- we store entries in nodes
- valid entries are non-NULL

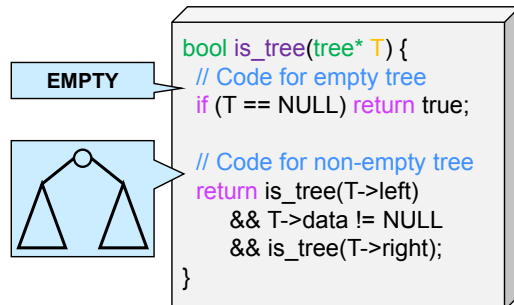


- This is a **recursive** function

- the **base case** is about the empty tree
- the **recursive case** is about every tree that is not empty
 - with a root
 - and two subtrees

A Minimal Tree Invariant


- We just check that the data field is never NULL



- But trees have constraints on their **structure**

- a node does not point to an ancestor
- a node has at most one parent

How to check them
is left as an exercise




- *What additional constraints on contents do we need to use trees to implement dictionaries?*

Binary Search Trees

Binary Search Trees


- *What additional constraints on the contents do we need to use trees to implement dictionaries?*
- Because lookup emulates binary search, the data in the tree need to be **ordered**
 - smaller values on the left
 - bigger values on the right



- A tree whose nodes are ordered is called a **binary search tree**

The BST Invariant

- A tree whose nodes are ordered is called a **binary search tree**



- We can write a specification function that check BSTs

